

HW3-Bhavik-Upadhyay

October 19, 2023

```
[1]: from typing import Optional, Callable
from typing_extensions import TypeAlias

import pandas as pd
import numpy as np
import nltk
import re
from bs4 import BeautifulSoup
import contractions

import os
from urllib import request
import gzip

from sklearn.model_selection import train_test_split
from sklearn.metrics import precision_score, recall_score, f1_score, \
    accuracy_score

from sklearn.linear_model import Perceptron, LogisticRegression
from sklearn.svm import LinearSVC
from sklearn.naive_bayes import MultinomialNB
```

```
[2]: import torch
import random

seed = 42
torch.manual_seed(seed)
random.seed(seed)
np.random.seed(0)
```

1 Task 1: Dataset Generation

```
[3]: url = 'https://web.archive.org/web/20201127142707if_/https://s3.amazonaws.com/
    amazon-reviews-pds/tsv/amazon_reviews_us_Office_Products_v1_00.tsv.gz'

extracted_file = 'data.tsv'
compressed_file = extracted_file + '.gz'
```

```

# Retrieve the dataset from given url and store it in location specified by
↳ compressed_file
if not os.path.exists(extracted_file):
    request.urlretrieve(url, compressed_file)

    # extract the dataset from the gzipped file
    with gzip.open(compressed_file, 'rb') as f_in, open(extracted_file, 'wb')
↳ as f_out:
        for line in f_in:
            f_out.write(line)

    os.remove(compressed_file)

# read the extracted data into pandas dataframe
original_df = pd.read_csv(extracted_file, sep='\t', on_bad_lines='skip',
↳ low_memory=False)
print(original_df.head())

```

	marketplace	customer_id	review_id	product_id	product_parent	\
0	US	43081963	R18RVCKGH1SSI9	B001BM2MAC	307809868	
1	US	10951564	R3L4L6LW1PUOFY	B00DZYEXPQ	75004341	
2	US	21143145	R2J8AWXWTDX2TF	B00RTMUHDW	529689027	
3	US	52782374	R1PR37BR7G3M6A	B00D7H8XB6	868449945	
4	US	24045652	R3BDDZMZBZDPU	B001XCWP34	33521401	

	product_title	product_category	\
0	Scotch Cushion Wrap 7961, 12 Inches x 100 Feet	Office Products	
1	Dust-Off Compressed Gas Duster, Pack of 4	Office Products	
2	Amram Tagger Standard Tag Attaching Tagging Gu...	Office Products	
3	AmazonBasics 12-Sheet High-Security Micro-Cut ...	Office Products	
4	Derwent Colored Pencils, Inktense Ink Pencils,...	Office Products	

	star_rating	helpful_votes	total_votes	vine	verified_purchase	\
0	5	0.0	0.0	N	Y	
1	5	0.0	1.0	N	Y	
2	5	0.0	0.0	N	Y	
3	1	2.0	3.0	N	Y	
4	4	0.0	0.0	N	Y	

	review_headline	\
0	Five Stars	
1	Phfffffft, Phfffffft. Lots of air, and it's C...	
2	but I am sure I will like it.	
3	and the shredder was dirty and the bin was par...	

4

Four Stars

	review_body	review_date
0	Great product.	2015-08-31
1	What's to say about this commodity item except...	2015-08-31
2	Haven't used yet, but I am sure I will like it.	2015-08-31
3	Although this was labeled as "new" the...	2015-08-31
4	Gorgeous colors and easy to use	2015-08-31

```
[4]: # creating the dataframe by taking only review_body and star_rating columns
df = pd.DataFrame(original_df[['review_body', 'star_rating']])
print(df.head())

# we notice there are some erroneous values for the star_rating column
print(df['star_rating'].unique())

# converting the star_rating to numeric values and dropping erroneous columns
df['star_rating'] = pd.to_numeric(df['star_rating'], errors='coerce')
df.dropna(inplace=True)

print(df['star_rating'].unique())
```

	review_body	star_rating
0	Great product.	5
1	What's to say about this commodity item except...	5
2	Haven't used yet, but I am sure I will like it.	5
3	Although this was labeled as "new" the...	1
4	Gorgeous colors and easy to use	4

['5' '1' '4' '2' '3' '2015-06-05' '2015-02-11' nan '2014-02-14']
[5. 1. 4. 2. 3.]

```
[5]: # creating the target column: target = 1 if star_rating is 1, 2 or 3. target = 2
      ↪ if star_rating is 4 or 5
df['star_rating'] = df['star_rating'].astype(int)
df['target'] = df['star_rating'].apply(lambda x: 0 if x <= 3 else 1)

sample_size = 50000

# creating a sample dataframe where target = 1 of size 50000 rows
class_1 = df.loc[df['target'] == 0].sample(n=sample_size, random_state=42)

# creating a sample dataframe where target = 2 of size 50000 rows
class_2 = df.loc[df['target'] == 1].sample(n=sample_size, random_state=42)

# merging the two sample dataframes
df_new = pd.concat([class_1, class_2], ignore_index=True)
```

```
[6]: def clean(review):
    """
    convert to lower-case
    remove html and urls
    remove non-alphabetical character
    remove extra spaces
    """

    # converting to lowercase
    review = review.lower()

    # removing htmls
    soup = BeautifulSoup(review, "html.parser")

    for a_tag in soup.find_all("a"):
        a_tag.decompose()

    review = soup.get_text()

    # removing urls
    review = re.sub(r'^https?:\/\/\.[^\s]*$', '', review)

    # removing non-alphabetical characters
    review = re.sub(r'[^a-zA-Z\s]', '', review)

    # removing extra spaces
    review = re.sub(r'\s+', ' ', review).strip()

    return review

df_new['review_body'] = df_new['review_body'].apply(clean)
```

```
C:\Users\bhavi\AppData\Local\Temp\ipykernel_17012\1465821108.py:13:
MarkupResemblesLocatorWarning: The input looks more like a filename than markup.
You may want to open this file and pass the filehandle into BeautifulSoup.
    soup = BeautifulSoup(review, "html.parser")
C:\Users\bhavi\AppData\Local\Temp\ipykernel_17012\1465821108.py:13:
MarkupResemblesLocatorWarning: The input looks more like a URL than markup. You
may want to use an HTTP client like requests to get the document behind the URL,
and feed that document to BeautifulSoup.
    soup = BeautifulSoup(review, "html.parser")
```

2 Task 2: Creating the Word2Vec Models

When comparing a pretrained model with a custom word2vec model, we find a substantial difference in vocabulary size: 3,000,000 unique words in the pretrained model versus 14,994 in the custom

model. This suggests the pretrained model may handle out-of-vocabulary words better during testing.

In semantic similarity tests: 1. Outstanding and excellent show lower similarity in the pretrained model. 2. The arithmetic “King - Man + Woman” correctly yields “Queen” in the pretrained model, but not in the custom model. 3. The arithmetic “Doctor - Man + Woman” doesn’t produce nurse-related results in the custom model.

Overall, the pretrained model performs better in some semantic tasks but struggles with similarity in specific cases.

2.1 Word2Vec from pretrained

```
[7]: import gensim.downloader
from gensim.models import KeyedVectors

# downloading the pre-trained model if it is not available
if not os.path.exists('pretrained_w2v.model'):
    pretrained_w2v = gensim.downloader.load('word2vec-google-news-300')
    pretrained_w2v.save('pretrained_w2v.model')
# if available, we load the data from local storage
else:
    pretrained_w2v = KeyedVectors.load('pretrained_w2v.model')
```

```
[8]: print(len(pretrained_w2v), len(pretrained_w2v[0]))
```

3000000 300

```
[9]: # printing the similarity score for outstanding and excellent
print(pretrained_w2v.similarity('outstanding', 'excellent'))
```

0.55674857

```
[10]: # printing the most similar words matching the arithmetic king - man + woman
print(pretrained_w2v.most_similar(positive=['king', 'woman'], negative=['man'],
    ↪topn=5))
```

[('queen', 0.7118193507194519), ('monarch', 0.6189674139022827), ('princess', 0.5902431011199951), ('crown_prince', 0.5499460697174072), ('prince', 0.5377321839332581)]

```
[11]: # printing the most similar words matching the arithmetic doctor - man + queen
print(pretrained_w2v.most_similar(positive=['doctor', 'woman'],
    ↪negative=['man'], topn=5))
```

[('gynecologist', 0.7093892097473145), ('nurse', 0.6477287411689758), ('doctors', 0.6471460461616516), ('physician', 0.6438996195793152), ('pediatrician', 0.6249487996101379)]

```
[12]: # Creating the vocabulary from the pretrained w2v
vocab = list(pretrained_w2v.index_to_key)

# Initialize an embedding matrix with zeros
embedding_dim = pretrained_w2v.vector_size
embedding_matrix = np.zeros((len(vocab), embedding_dim), dtype=np.float32)

# Populate the embedding matrix with Word2Vec vectors
for i, word in enumerate(vocab):
    if word in pretrained_w2v:
        # print(word, i, custom_w2v.wv[word])
        embedding_matrix[i] = pretrained_w2v[word]

print(len(embedding_matrix), len(embedding_matrix[0]))

3000000 300
```

2.2 From dataset (custom)

```
[13]: from gensim.models import Word2Vec

# getting the tokens
tokenized = [nltk.word_tokenize(review) for review in df_new['review_body']]

# if the model is not already available, then create from scratch
if not os.path.exists('custom_w2v.model'):
    custom_w2v = Word2Vec(tokenized, vector_size=300, window=13, min_count=9,
        ↪sg=1, workers=1)
    custom_w2v.save('custom_w2v.model')
# if available, load the model from local storage
else:
    custom_w2v = Word2Vec.load('custom_w2v.model')
```

```
[14]: print(len(custom_w2v.wv), len(custom_w2v.wv[0]))
```

14994 300

```
[15]: # printing the similarity score for outstanding and excellent
print(custom_w2v.wv.similarity('outstanding', 'excellent'))
```

0.6205609

```
[16]: # printing the most similar words matching the arithmetic king - man + woman
print(custom_w2v.wv.most_similar(positive=['king', 'woman'], negative=['man'],
    ↪topn=5))
```

```
[('Idea', 0.5137808918952942), ('Archival', 0.5026917457580566), ('fine-point',
0.49942412972450256), ('inherited', 0.4969730079174042), ('Flair',
0.4938381314277649)]
```

```
[17]: # printing the most similar words matching the arithmetic doctor - man + queen
print(custom_w2v.wv.most_similar(positive=['doctor', 'woman'],
↪negative=['man'], topn=5))
```

```
[('assistant', 0.4972769320011139), ('visiting', 0.4651092290878296),
('deadlines', 0.45854732394218445), ('appointment', 0.45726412534713745),
('prayer', 0.4571249485015869)]
```

3 Task 3: Simple Models

- **Perceptron** achieved **80.77% accuracy** on word embeddings, which is slightly better than **79.345%** using TF-IDF.
- **SVM** attained **82.665% accuracy** on word embeddings, slightly lower than **84.865%** using TF-IDF.
- In general, word embeddings exhibit competitive performance, however, in the SVM task, we notice slightly less performance.

Note: The two figures below show the accuracy obtained for perceptron and SVM when training on TF-IDF. These were obtained by modifying the metric of calculation in Homework 1 without any changes as to how the models were trained or features were extracted.

```
In [12]: 1 # Creating and training the perceptron for bag-of-words
2 bow_clf = Perceptron(penalty='elasticnet', l1_ratio=0.1, eta0=1e-3, alpha=1e-6, tol=1e-4, random_state=42)
3 bow_clf.fit(bow_X_train, Y_train)
4
5 # making predictions on the bag-of-words test set
6 bow_Y_pred = bow_clf.predict(bow_X_test)
7
8 print('BOW: ', accuracy_score(Y_test, bow_Y_pred))
9
10 # Creating and training the perceptron for TF-IDF
11 tf_idf_clf = Perceptron(penalty='elasticnet', l1_ratio=0.3, eta0=1e-5, max_iter=1000, alpha=1e-6, tol=1e-4, random_state=42)
12 tf_idf_clf.fit(tf_idf_X_train, Y_train)
13
14 # making predictions on the TF-IDF test set
15 tf_idf_Y_pred = tf_idf_clf.predict(tf_idf_X_test)
16
17 print('TF-IDF: ', accuracy_score(Y_test, tf_idf_Y_pred))

BOW: 0.79925
TF-IDF: 0.79345
```

Accuracy of perceptron on TF-IDF

```
In [14]: 1 # Creating and training SVM model on Bag-of-words training set
2 bow_svm = LinearSVC(max_iter=1000, penalty='l1', dual=False, C=0.1, random_state=42)
3 bow_svm.fit(bow_X_train, Y_train)
4
5 # making predictions on bag-of-words test set
6 bow_Y_pred = bow_svm.predict(bow_X_test)
7
8 print('BOW:', accuracy_score(Y_test, bow_Y_pred))
9
10
11 # Creating and training SVM model on TFIDF training set
12 tf_idf_svm = LinearSVC(max_iter=10, dual=False, C=0.1, random_state=42)
13 tf_idf_svm.fit(tf_idf_X_train, Y_train)
14
15 # making predictions on TF-IDF test set
16 tf_idf_Y_pred = tf_idf_svm.predict(tf_idf_X_test)
17
18 print('TF-IDF: ', accuracy_score(Y_test, tf_idf_Y_pred))

BOW: 0.8429
TF-IDF: 0.84865
```

Accuracy of SVM on TF-IDF

3.0.1 Creating mean sentence embeddings

```
[18]: # a function to create mean embeddings for a sentence given a w2v model
def create_avg_embeddings(sentence, w2v):
    tokens = nltk.word_tokenize(sentence)
    vectors = [w2v[word] for word in tokens if word in w2v]

    if vectors:
        embedding = np.mean(vectors, axis=0, dtype=np.float32)
    else:
        embedding = np.zeros(w2v.vector_size, dtype=np.float32)

    return embedding

[19]: avg_embeddings = df_new['review_body'].apply(lambda x: create_avg_embeddings(x,
↳ pretrained_w2v))
avg_embeddings = np.array(avg_embeddings.tolist())

targets = df_new['target']

[20]: X_train_avg, X_test_avg, Y_train_avg, Y_test_avg = train_test_split(
    avg_embeddings,
    df_new['target'],
    shuffle=True,
    test_size=0.2,
    random_state=42
)

Y_train_avg = np.array(Y_train_avg.tolist())
Y_test_avg = np.array(Y_test_avg.tolist())
```

3.0.2 Perceptron training

```
[21]: per_clf = Perceptron(penalty='elasticnet', l1_ratio=0.8, alpha=1e-5, tol=1e-4,
↳ random_state=42)
per_clf.fit(list(X_train_avg), Y_train_avg)

per_Y_preds = per_clf.predict(list(X_test_avg))
per_acc = accuracy_score(per_Y_preds, Y_test_avg)
print(per_acc)
```

0.8077

3.0.3 SVM training

```
[22]: svc_clf = LinearSVC(dual=True, loss='hinge', C=0.8, max_iter=10000,
    ↪random_state=42)
svc_clf.fit(list(X_train_avg), Y_train_avg)

svc_Y_preds = svc_clf.predict(list(X_test_avg))
svc_acc = accuracy_score(svc_Y_preds, Y_test_avg)
print(svc_acc)
```

0.82665

4 Task 4: Feedforward Neural Network

```
[23]: # general purpose function for training a model with given training data
def train(model, train_data, test_data, criterion, optimizer, n_epochs=10,
    ↪verbose=True, device='cpu'):
    for epoch in range(n_epochs):
        train_loss = 0.
        test_loss = 0.

        model.train()
        for data, target in train_data:
            # shifting data and target to the appropriate device
            data = data.to(device)
            target = target.to(device)

            # setting the gradients to zero
            optimizer.zero_grad()

            # getting the output and calculating the loss
            out = model(data)
            loss = criterion(out, target)

            # performing the backward step and using the optimizer
            loss.backward()
            optimizer.step()

            train_loss += loss.item()

        model.eval()
        with torch.no_grad():
            for data, target in test_data:
                # shifting data and target to appropriate device
                data = data.to(device)
                target = target.to(device)
```

```

        # getting the output and then getting the loss and updating the
↪total test loss
        out = model(data)

        loss = criterion(out, target)
        test_loss += loss.item()

    if verbose:
        print(f'Epoch: {epoch+1} / {n_epochs}\tTraining Loss:
↪{train_loss}\tTest Loss: {test_loss}')

```

```

[24]: # a function to calculate accuracy on test data for a given model
def accuracy(model, test_data, device='cpu'):
    correct, total = 0, 0

    with torch.no_grad():
        for data, target in test_data:
            # shifting data and target to appropriate device
            data = data.to(device)
            target = target.to(device)

            # getting the output and predictions
            out = model(data)
            _, preds = torch.max(out.data, 1)

            # updating the total and correct variables
            total += target.size(0)
            correct += (preds == target).sum().item()

    return (100*correct) / total

```

```

[25]: import torch
from torch.utils.data import TensorDataset, DataLoader
import torch.nn as nn
import torch.optim as optim

device = 'cuda:0' if torch.cuda.is_available() else 'cpu'

# converting numpy arrays created for perceptron and svm to torch tensors
X_train_avg = torch.tensor(X_train_avg, dtype=torch.float32)
X_test_avg = torch.tensor(X_test_avg, dtype=torch.float32)

Y_train_avg = torch.tensor(Y_train_avg, dtype=torch.int64)
Y_test_avg = torch.tensor(Y_test_avg, dtype=torch.int64)

# generating dataset from created tensors
train_dataset1 = TensorDataset(X_train_avg, Y_train_avg)

```

```

test_dataset1 = TensorDataset(X_test_avg, Y_test_avg)

batch_size = 256
# creating train and test data loaders
train_loader1 = DataLoader(train_dataset1, batch_size=batch_size, shuffle=True)
test_loader1 = DataLoader(test_dataset1, batch_size=batch_size)

print(device)

```

cuda:0

4.1 Task 4 (a)

- Here, we train a feedforward neural network on sentence embeddings obtained by calculating the mean of all word embeddings in the sentence.
- We train the neural network using AdamW optimizer with $1e-4$ learning rate for 100 epochs with batch size 256.
- The accuracy on test set for this model is roughly between **82% - 84%**.

```

[26]: class NeuralNetwork1(nn.Module):
    def __init__(self):
        super(NeuralNetwork1, self).__init__()
        self.embedding_dim = 300
        self.hidden1 = 50
        self.hidden2 = 5
        self.out_dim = 2

        self.linear = nn.Sequential(
            nn.Linear(self.embedding_dim, self.hidden1),
            nn.ReLU(),
            nn.Linear(self.hidden1, self.hidden2),
            nn.ReLU(),
            nn.Linear(self.hidden2, self.out_dim),

        )

    def forward(self, x):
        x = self.linear(x)

        return x

model1 = NeuralNetwork1()
model1.to(device)
print(model1)

```

```

NeuralNetwork1(
  (linear): Sequential(
    (0): Linear(in_features=300, out_features=50, bias=True)
    (1): ReLU()
  )
)

```

```

        (2): Linear(in_features=50, out_features=5, bias=True)
        (3): ReLU()
        (4): Linear(in_features=5, out_features=2, bias=True)
    )
)

```

```

[27]: criterion = nn.CrossEntropyLoss()
      optimizer = torch.optim.AdamW(model1.parameters(), lr=1e-4)

```

```

[28]: # training the model
      train(model1, train_loader1, test_loader1, criterion, optimizer, n_epochs=100,
      ↪device=device)

```

Epoch: 1 / 100	Training Loss: 223.27994138002396	Test Loss: 54.716657280921936
Epoch: 2 / 100	Training Loss: 205.70290875434875	Test Loss: 48.29630637168884
Epoch: 3 / 100	Training Loss: 176.52640929818153	Test Loss: 41.207524448633194
Epoch: 4 / 100	Training Loss: 155.0030519068241	Test Loss: 37.60107463598251
Epoch: 5 / 100	Training Loss: 144.89616572856903	Test Loss: 35.95973256230354
Epoch: 6 / 100	Training Loss: 140.06317156553268	Test Loss: 35.0801557302475
Epoch: 7 / 100	Training Loss: 137.1439170241356	Test Loss: 34.57513916492462
Epoch: 8 / 100	Training Loss: 135.18273961544037	Test Loss: 34.08438017964363
Epoch: 9 / 100	Training Loss: 133.62331506609917	Test Loss: 33.74214479327202
Epoch: 10 / 100	Training Loss: 132.32582929730415	Test Loss: 33.46182382106781
Epoch: 11 / 100	Training Loss: 131.22946453094482	Test Loss: 33.21343466639519
Epoch: 12 / 100	Training Loss: 130.28965264558792	Test Loss: 32.998418152332306
Epoch: 13 / 100	Training Loss: 129.46428593993187	Test Loss: 32.80886098742485
Epoch: 14 / 100	Training Loss: 128.77806401252747	Test Loss: 32.64651048183441
Epoch: 15 / 100	Training Loss: 128.14307433366776	Test Loss: 32.497187316417694
Epoch: 16 / 100	Training Loss: 127.5267682671547	Test Loss: 32.354840099811554
Epoch: 17 / 100	Training Loss: 126.88095933198929	Test Loss: 32.239379823207855
Epoch: 18 / 100	Training Loss: 126.27783998847008	Test Loss:

32.22042775154114	
Epoch: 19 / 100 Training Loss: 125.91750055551529	Test Loss:
31.992936491966248	
Epoch: 20 / 100 Training Loss: 125.32829281687737	Test Loss:
31.872982531785965	
Epoch: 21 / 100 Training Loss: 124.9611741900444	Test Loss:
31.765440344810486	
Epoch: 22 / 100 Training Loss: 124.49292731285095	Test Loss:
31.662551641464233	
Epoch: 23 / 100 Training Loss: 124.03318306803703	Test Loss:
31.565906643867493	
Epoch: 24 / 100 Training Loss: 123.64501696825027	Test Loss:
31.490920424461365	
Epoch: 25 / 100 Training Loss: 123.26009619235992	Test Loss:
31.418159753084183	
Epoch: 26 / 100 Training Loss: 122.91045781970024	Test Loss:
31.33314400911331	
Epoch: 27 / 100 Training Loss: 122.6216288805008	Test Loss:
31.26932805776596	
Epoch: 28 / 100 Training Loss: 122.26745873689651	Test Loss:
31.225706696510315	
Epoch: 29 / 100 Training Loss: 121.92679944634438	Test Loss:
31.129334658384323	
Epoch: 30 / 100 Training Loss: 121.6411349773407	Test Loss:
31.135849595069885	
Epoch: 31 / 100 Training Loss: 121.3110009431839	Test Loss:
31.019828230142593	
Epoch: 32 / 100 Training Loss: 121.060700237751	Test Loss: 30.958229959011078
Epoch: 33 / 100 Training Loss: 120.7994986474514	Test Loss:
30.945522993803024	
Epoch: 34 / 100 Training Loss: 120.50094133615494	Test Loss:
30.85480245947838	
Epoch: 35 / 100 Training Loss: 120.2346733212471	Test Loss:
30.826304495334625	
Epoch: 36 / 100 Training Loss: 120.0566368997097	Test Loss:
30.808255553245544	
Epoch: 37 / 100 Training Loss: 119.85122066736221	Test Loss:
30.851676404476166	
Epoch: 38 / 100 Training Loss: 119.64552560448647	Test Loss:
30.70193523168564	
Epoch: 39 / 100 Training Loss: 119.43237388134003	Test Loss:
30.66961708664894	
Epoch: 40 / 100 Training Loss: 119.18115195631981	Test Loss:
30.64968267083168	
Epoch: 41 / 100 Training Loss: 119.03138810396194	Test Loss:
30.572869837284088	
Epoch: 42 / 100 Training Loss: 118.7706771492958	Test Loss:
30.539514303207397	

Epoch: 43 / 100 Training Loss: 118.66438245773315 30.56415206193924	Test Loss:
Epoch: 44 / 100 Training Loss: 118.44496589899063 30.49730333685875	Test Loss:
Epoch: 45 / 100 Training Loss: 118.27230963110924 30.445301681756973	Test Loss:
Epoch: 46 / 100 Training Loss: 118.13230195641518 30.425574600696564	Test Loss:
Epoch: 47 / 100 Training Loss: 117.95960614085197 30.442767202854156	Test Loss:
Epoch: 48 / 100 Training Loss: 117.84438782930374 30.37128135561943	Test Loss:
Epoch: 49 / 100 Training Loss: 117.6802129149437 30.35201469063759	Test Loss:
Epoch: 50 / 100 Training Loss: 117.48000007867813 30.369530647993088	Test Loss:
Epoch: 51 / 100 Training Loss: 117.35213413834572 30.313990354537964	Test Loss:
Epoch: 52 / 100 Training Loss: 117.1872521340847 30.272311061620712	Test Loss:
Epoch: 53 / 100 Training Loss: 117.05581396818161 30.287183433771133	Test Loss:
Epoch: 54 / 100 Training Loss: 116.89013743400574 30.219446301460266	Test Loss:
Epoch: 55 / 100 Training Loss: 116.75635251402855 30.19229105114937	Test Loss:
Epoch: 56 / 100 Training Loss: 116.64956653118134 30.226687908172607	Test Loss:
Epoch: 57 / 100 Training Loss: 116.44064235687256 30.174104303121567	Test Loss:
Epoch: 58 / 100 Training Loss: 116.2725133895874 30.144307047128677	Test Loss:
Epoch: 59 / 100 Training Loss: 116.18694359064102 30.14928701519966	Test Loss:
Epoch: 60 / 100 Training Loss: 116.06538730859756 30.069722801446915	Test Loss:
Epoch: 61 / 100 Training Loss: 115.90143203735352 30.04452556371689	Test Loss:
Epoch: 62 / 100 Training Loss: 115.8383446931839 30.03296783566475	Test Loss:
Epoch: 63 / 100 Training Loss: 115.65640524029732 30.038346081972122	Test Loss:
Epoch: 64 / 100 Training Loss: 115.46543255448341 29.980317026376724	Test Loss:
Epoch: 65 / 100 Training Loss: 115.36341765522957 29.972163796424866	Test Loss:
Epoch: 66 / 100 Training Loss: 115.20414933562279 29.93922859430313	Test Loss:

Epoch: 67 / 100 Training Loss: 115.07365503907204	Test Loss:
29.927978098392487	
Epoch: 68 / 100 Training Loss: 114.92494955658913	Test Loss:
29.88445019721985	
Epoch: 69 / 100 Training Loss: 114.82510590553284	Test Loss:
29.96846315264702	
Epoch: 70 / 100 Training Loss: 114.70393919944763	Test Loss:
29.846114546060562	
Epoch: 71 / 100 Training Loss: 114.50519832968712	Test Loss:
29.823754519224167	
Epoch: 72 / 100 Training Loss: 114.43318873643875	Test Loss:
29.961664497852325	
Epoch: 73 / 100 Training Loss: 114.29965949058533	Test Loss:
29.841556757688522	
Epoch: 74 / 100 Training Loss: 114.14315912127495	Test Loss:
29.775625854730606	
Epoch: 75 / 100 Training Loss: 114.02090355753899	Test Loss:
29.751518547534943	
Epoch: 76 / 100 Training Loss: 113.9025110900402	Test Loss:
29.726153671741486	
Epoch: 77 / 100 Training Loss: 113.76714563369751	Test Loss:
29.707972019910812	
Epoch: 78 / 100 Training Loss: 113.66112577915192	Test Loss:
29.693700343370438	
Epoch: 79 / 100 Training Loss: 113.52604535222054	Test Loss:
29.68906545639038	
Epoch: 80 / 100 Training Loss: 113.44295200705528	Test Loss:
29.694132387638092	
Epoch: 81 / 100 Training Loss: 113.22596323490143	Test Loss:
29.699741423130035	
Epoch: 82 / 100 Training Loss: 113.16738465428352	Test Loss:
29.632641434669495	
Epoch: 83 / 100 Training Loss: 113.02845710515976	Test Loss:
29.603986263275146	
Epoch: 84 / 100 Training Loss: 112.89375838637352	Test Loss:
29.577137231826782	
Epoch: 85 / 100 Training Loss: 112.86698698997498	Test Loss:
29.59031268954277	
Epoch: 86 / 100 Training Loss: 112.68282690644264	Test Loss:
29.568561047315598	
Epoch: 87 / 100 Training Loss: 112.57322052121162	Test Loss:
29.532892733812332	
Epoch: 88 / 100 Training Loss: 112.46350705623627	Test Loss:
29.514625161886215	
Epoch: 89 / 100 Training Loss: 112.35060584545135	Test Loss:
29.497875690460205	
Epoch: 90 / 100 Training Loss: 112.22641348838806	Test Loss:
29.4856516122818	

Epoch: 91 / 100 Training Loss: 112.12844929099083	Test Loss:
29.459724247455597	
Epoch: 92 / 100 Training Loss: 111.98545953631401	Test Loss:
29.45690569281578	
Epoch: 93 / 100 Training Loss: 111.86591139435768	Test Loss:
29.447207391262054	
Epoch: 94 / 100 Training Loss: 111.7561075091362	Test Loss:
29.399598449468613	
Epoch: 95 / 100 Training Loss: 111.67023974657059	Test Loss:
29.484870731830597	
Epoch: 96 / 100 Training Loss: 111.57436427474022	Test Loss:
29.38424116373062	
Epoch: 97 / 100 Training Loss: 111.46817779541016	Test Loss:
29.367492109537125	
Epoch: 98 / 100 Training Loss: 111.29240453243256	Test Loss:
29.386514335870743	
Epoch: 99 / 100 Training Loss: 111.14512953162193	Test Loss:
29.411173075437546	
Epoch: 100 / 100 Training Loss: 111.15172135829926	Test Loss:
29.344766169786453	

```
[29]: part_4a_accuracy = accuracy(model1, test_loader1, device)
      print(part_4a_accuracy)
```

83.965

```
[30]: def create_pad_embeddings(sentence, w2v, max_len=10):
      tokens = nltk.word_tokenize(sentence)
      vec_size = w2v.vector_size

      embedding = np.zeros((max_len * vec_size), dtype=np.float32)
      for i, word in enumerate(tokens):
          if i >= max_len:
              break

          if word in w2v:
              embedding[i*vec_size: (i+1)*vec_size] = w2v[word]

      return embedding

      sentence = "This is a sample sentence"
      sample_embedding = create_pad_embeddings(sentence, pretrained_w2v)
      print(sample_embedding.shape)

      print(create_avg_embeddings(sentence, pretrained_w2v).shape)
```

(3000,)

(300,)


```
[31]: pad_embeddings = df_new['review_body'].apply(lambda x: create_pad_embeddings(x,
↳ pretrained_w2v))
pad_embeddings = np.array(pad_embeddings.tolist())

[32]: X_train_pad, X_test_pad, Y_train_pad, Y_test_pad = train_test_split(
    pad_embeddings,
    targets,
    shuffle=True,
    test_size=0.2,
    random_state=42
)

Y_train_pad = np.array(Y_train_pad.tolist())
Y_test_pad = np.array(Y_test_pad.tolist())

[33]: X_train_pad = torch.tensor(X_train_pad, dtype=torch.float32)
X_test_pad = torch.tensor(X_test_pad, dtype=torch.float32)

Y_train_pad = torch.tensor(Y_train_pad, dtype=torch.int64)
Y_test_pad = torch.tensor(Y_test_pad, dtype=torch.int64)

train_dataset2 = TensorDataset(X_train_pad, Y_train_pad)
test_dataset2 = TensorDataset(X_test_pad, Y_test_pad)

batch_size = 256
train_loader2 = DataLoader(train_dataset2, batch_size=batch_size, shuffle=True)
test_loader2 = DataLoader(test_dataset2, batch_size=batch_size)
```

4.2 Task 4 (b)

- Here, we train a feedforward neural network on sentence embeddings obtained by concatenating the first 10 word embeddings and applying padding if the sentence is smaller than 10 words.
- We use a ReLU activation layer between the linear layers.
- We train the neural network using AdamW optimizer with 1e-4 learning rate and 1e-4 weight decay for 30 epochs with batch size 256.
- Upon re-running the code several times, we observe the accuracy on test set for this model is roughly between **75.5% - 76.6%**.

```
[34]: class NeuralNetwork2(nn.Module):
    def __init__(self):
        super(NeuralNetwork2, self).__init__()
        self.embedding_dim = 3000
        self.hidden1 = 50
        self.hidden2 = 5
        self.out_dim = 2
```

```

        self.linear = nn.Sequential(
            nn.Linear(self.embedding_dim, self.hidden1),
            nn.ReLU(),
            nn.Linear(self.hidden1, self.hidden2),
            nn.ReLU(),
            nn.Linear(self.hidden2, self.out_dim),
        )

    def forward(self, x):
        x = self.linear(x)

        return x

model2 = NeuralNetwork2()
model2.to(device)
print(model2)

```

```

NeuralNetwork2(
  (linear): Sequential(
    (0): Linear(in_features=3000, out_features=50, bias=True)
    (1): ReLU()
    (2): Linear(in_features=50, out_features=5, bias=True)
    (3): ReLU()
    (4): Linear(in_features=5, out_features=2, bias=True)
  )
)

```

```
[35]: optimizer = torch.optim.AdamW(model2.parameters(), lr=1e-4, weight_decay=1e-4)
```

```
[36]: train(model2, train_loader2, test_loader2, criterion=criterion,
        ↪optimizer=optimizer, n_epochs=30, device=device)
```

Epoch: 1 / 30	Training Loss: 194.17466282844543	Test Loss: 43.72030174732208
Epoch: 2 / 30	Training Loss: 165.8945385813713	Test Loss: 41.193147748708725
Epoch: 3 / 30	Training Loss: 159.00936275720596	Test Loss: 40.39856415987015
Epoch: 4 / 30	Training Loss: 155.71505045890808	Test Loss: 39.918146044015884
Epoch: 5 / 30	Training Loss: 153.38784858584404	Test Loss: 39.65698781609535
Epoch: 6 / 30	Training Loss: 151.5978156030178	Test Loss: 39.41517451405525
Epoch: 7 / 30	Training Loss: 149.9291484951973	Test Loss: 39.31141784787178
Epoch: 8 / 30	Training Loss: 148.38147443532944	Test Loss: 39.11142221093178

Epoch: 9 / 30	Training Loss: 146.90436762571335	Test Loss:
	39.007114231586456	
Epoch: 10 / 30	Training Loss: 145.52796256542206	Test Loss:
	38.82473134994507	
Epoch: 11 / 30	Training Loss: 144.09228175878525	Test Loss:
	38.75255364179611	
Epoch: 12 / 30	Training Loss: 142.6417380273342	Test Loss:
	38.639875173568726	
Epoch: 13 / 30	Training Loss: 141.23964083194733	Test Loss:
	38.59825983643532	
Epoch: 14 / 30	Training Loss: 140.0570511519909	Test Loss:
	38.49330136179924	
Epoch: 15 / 30	Training Loss: 138.561135917902	Test Loss: 38.5560596883297
Epoch: 16 / 30	Training Loss: 137.25511133670807	Test Loss:
	38.54057967662811	
Epoch: 17 / 30	Training Loss: 135.8898992240429	Test Loss:
	38.45652109384537	
Epoch: 18 / 30	Training Loss: 134.39312362670898	Test Loss:
	38.396814465522766	
Epoch: 19 / 30	Training Loss: 133.03255796432495	Test Loss:
	38.37224414944649	
Epoch: 20 / 30	Training Loss: 131.66759631037712	Test Loss:
	38.39131438732147	
Epoch: 21 / 30	Training Loss: 130.14192607998848	Test Loss:
	38.43498423695564	
Epoch: 22 / 30	Training Loss: 128.69674035906792	Test Loss:
	38.42015391588211	
Epoch: 23 / 30	Training Loss: 127.28686335682869	Test Loss:
	38.41508102416992	
Epoch: 24 / 30	Training Loss: 125.83272141218185	Test Loss:
	38.49851316213608	
Epoch: 25 / 30	Training Loss: 124.14575991034508	Test Loss:
	38.61513492465019	
Epoch: 26 / 30	Training Loss: 122.76370960474014	Test Loss:
	38.81642150878906	
Epoch: 27 / 30	Training Loss: 121.15855580568314	Test Loss:
	38.708027839660645	
Epoch: 28 / 30	Training Loss: 119.53005722165108	Test Loss:
	38.78369262814522	
Epoch: 29 / 30	Training Loss: 117.86586755514145	Test Loss:
	38.875079065561295	
Epoch: 30 / 30	Training Loss: 116.18799751996994	Test Loss:
	38.919744431972504	

```
[37]: part_4b_accuracy = accuracy(model2, test_loader2, device=device)
      print(part_4b_accuracy)
```

76.29

5 Task 5: Recurrent Neural Networks

```
[38]: def create_block_embeddings(sentence, w2v, max_len=10):
    tokens = nltk.word_tokenize(sentence)
    vec_size = w2v.vector_size

    embedding = np.zeros((max_len, vec_size), dtype=np.float32)
    for i, word in enumerate(tokens):
        if i >= max_len:
            break

        if word in w2v:
            embedding[i] = w2v[word]

    return embedding

sentence = "This is a sample sentence"
sample_embedding = create_block_embeddings(sentence, pretrained_w2v)
```

```
[39]: block_embeddings = df_new['review_body'].apply(lambda x:
    ↪ create_block_embeddings(x, pretrained_w2v))
block_embeddings = np.array(block_embeddings.tolist())
```

```
[40]: X_train_block, X_test_block, Y_train_block, Y_test_block = train_test_split(
    block_embeddings,
    targets,
    shuffle=True,
    test_size=0.2,
    random_state=42
)

Y_train_block = np.array(Y_train_block.tolist())
Y_test_block = np.array(Y_test_block.tolist())
```

```
[41]: X_train_block = torch.tensor(X_train_block, dtype=torch.float32)
X_test_block = torch.tensor(X_test_block, dtype=torch.float32)

Y_train_block = torch.tensor(Y_train_block, dtype=torch.int64)
Y_test_block = torch.tensor(Y_test_block, dtype=torch.int64)

train_dataset3 = TensorDataset(X_train_block, Y_train_block)
test_dataset3 = TensorDataset(X_test_block, Y_test_block)

batch_size = 256
train_loader3 = DataLoader(train_dataset3, batch_size=batch_size, shuffle=True)
test_loader3 = DataLoader(test_dataset3, batch_size=batch_size)
```

5.1 Task 5 (a)

- For this task, we train a RNN on sentence embeddings obtained by stacking the first 10 word embeddings and applying padding if the sentence is smaller than 10 words.
- We train the neural network using AdamW optimizer with 1e-3 learning rate with 1e-6 weight decay for 100 epochs with batch size 256.
- Upon re-running the code several times, we observe the accuracy on test set for this model is roughly around **77.7% - 78.5%**.

```
[42]: class RNNClf(nn.Module):
    def __init__(self):
        super(RNNClf, self).__init__()
        self.embed_size = 300
        self.hidden_size = 10
        self.out_size = 2

        self.rnn = nn.RNN(input_size=self.embed_size, hidden_size=self.
↪hidden_size, batch_first=True)
        self.linear = nn.Linear(self.hidden_size, self.out_size)

    def forward(self, x):
        x, hidden = self.rnn(x)
        x = self.linear(x[:, -1, :])

        return x

rnnModel = RNNClf()
rnnModel.to(device)
print(rnnModel)
```

```
RNNClf(
  (rnn): RNN(300, 10, batch_first=True)
  (linear): Linear(in_features=10, out_features=2, bias=True)
)
```

```
[43]: optimizer = torch.optim.AdamW(rnnModel.parameters(), lr=1e-3, weight_decay=1e-6)
```

```
[44]: train(rnnModel, train_loader3, test_loader3, criterion, optimizer,
↪n_epochs=100, device=device)
```

Epoch: 1 / 100	Training Loss: 189.24623772501945	Test Loss: 42.86137253046036
Epoch: 2 / 100	Training Loss: 165.2519319653511	Test Loss: 41.2786665558815
Epoch: 3 / 100	Training Loss: 160.46886545419693	Test Loss: 40.46919998526573
Epoch: 4 / 100	Training Loss: 158.11708521842957	Test Loss: 40.11415392160416
Epoch: 5 / 100	Training Loss: 156.1327583193779	Test Loss:

39.94943976402283	
Epoch: 6 / 100 Training Loss: 154.3838656246662	Test Loss:
39.207602590322495	
Epoch: 7 / 100 Training Loss: 153.1704162955284	Test Loss:
39.585414320230484	
Epoch: 8 / 100 Training Loss: 151.9054266512394	Test Loss:
38.63105762004852	
Epoch: 9 / 100 Training Loss: 150.69831427931786	Test Loss:
38.46751955151558	
Epoch: 10 / 100 Training Loss: 149.9881165921688	Test Loss:
38.16771939396858	
Epoch: 11 / 100 Training Loss: 149.20873829722404	Test Loss:
38.12427684664726	
Epoch: 12 / 100 Training Loss: 148.06970876455307	Test Loss:
37.88624459505081	
Epoch: 13 / 100 Training Loss: 147.25514441728592	Test Loss:
37.83639848232269	
Epoch: 14 / 100 Training Loss: 147.02438268065453	Test Loss:
37.81947907805443	
Epoch: 15 / 100 Training Loss: 146.1519265472889	Test Loss:
37.40324741601944	
Epoch: 16 / 100 Training Loss: 145.35491436719894	Test Loss:
37.36286148428917	
Epoch: 17 / 100 Training Loss: 144.6717321574688	Test Loss:
37.37463703751564	
Epoch: 18 / 100 Training Loss: 143.9930343925953	Test Loss:
37.337947338819504	
Epoch: 19 / 100 Training Loss: 143.5869961977005	Test Loss:
37.078449696302414	
Epoch: 20 / 100 Training Loss: 143.13337874412537	Test Loss:
37.605335503816605	
Epoch: 21 / 100 Training Loss: 142.94152796268463	Test Loss:
37.53921481966972	
Epoch: 22 / 100 Training Loss: 142.4502227306366	Test Loss:
36.98101082444191	
Epoch: 23 / 100 Training Loss: 141.7295179963112	Test Loss:
36.969017028808594	
Epoch: 24 / 100 Training Loss: 141.79342359304428	Test Loss:
36.92888504266739	
Epoch: 25 / 100 Training Loss: 141.24457421898842	Test Loss:
36.921512484550476	
Epoch: 26 / 100 Training Loss: 140.69035521149635	Test Loss:
36.67476940155029	
Epoch: 27 / 100 Training Loss: 140.56993076205254	Test Loss:
36.67093303799629	
Epoch: 28 / 100 Training Loss: 140.12270125746727	Test Loss:
36.673606127500534	
Epoch: 29 / 100 Training Loss: 140.30863592028618	Test Loss:

36.74337786436081	
Epoch: 30 / 100 Training Loss: 139.8257461488247	Test Loss:
36.717282712459564	
Epoch: 31 / 100 Training Loss: 139.38076090812683	Test Loss:
36.759630620479584	
Epoch: 32 / 100 Training Loss: 139.37510937452316	Test Loss:
36.75336068868637	
Epoch: 33 / 100 Training Loss: 138.59608009457588	Test Loss:
36.559115529060364	
Epoch: 34 / 100 Training Loss: 138.47659105062485	Test Loss:
36.69157314300537	
Epoch: 35 / 100 Training Loss: 138.35360470414162	Test Loss:
36.672885090112686	
Epoch: 36 / 100 Training Loss: 138.4716020822525	Test Loss:
36.34186860918999	
Epoch: 37 / 100 Training Loss: 138.0806143283844	Test Loss:
36.49436393380165	
Epoch: 38 / 100 Training Loss: 137.92330566048622	Test Loss:
36.363194674253464	
Epoch: 39 / 100 Training Loss: 137.7402704358101	Test Loss:
36.36631777882576	
Epoch: 40 / 100 Training Loss: 137.59592580795288	Test Loss:
36.46004328131676	
Epoch: 41 / 100 Training Loss: 137.2826405465603	Test Loss:
36.5595398247242	
Epoch: 42 / 100 Training Loss: 137.12348607182503	Test Loss:
36.4635514318943	
Epoch: 43 / 100 Training Loss: 137.01033291220665	Test Loss:
36.34114542603493	
Epoch: 44 / 100 Training Loss: 136.74183830618858	Test Loss:
36.264641135931015	
Epoch: 45 / 100 Training Loss: 136.3272634446621	Test Loss:
36.134812384843826	
Epoch: 46 / 100 Training Loss: 136.4916720688343	Test Loss:
36.68380865454674	
Epoch: 47 / 100 Training Loss: 136.05500841140747	Test Loss:
36.26983544230461	
Epoch: 48 / 100 Training Loss: 135.8873808979988	Test Loss:
36.39075103402138	
Epoch: 49 / 100 Training Loss: 136.05873787403107	Test Loss:
36.15823784470558	
Epoch: 50 / 100 Training Loss: 135.63928240537643	Test Loss:
36.31730031967163	
Epoch: 51 / 100 Training Loss: 136.06965699791908	Test Loss:
36.068804532289505	
Epoch: 52 / 100 Training Loss: 135.45133262872696	Test Loss:
36.76198589801788	
Epoch: 53 / 100 Training Loss: 135.2078354358673	Test Loss:

36.18658712506294	
Epoch: 54 / 100 Training Loss: 135.38859137892723	Test Loss:
36.101177006959915	
Epoch: 55 / 100 Training Loss: 135.46006244421005	Test Loss:
36.34166017174721	
Epoch: 56 / 100 Training Loss: 135.55552461743355	Test Loss:
35.95088368654251	
Epoch: 57 / 100 Training Loss: 135.10890033841133	Test Loss:
37.04864928126335	
Epoch: 58 / 100 Training Loss: 134.89186203479767	Test Loss:
36.29133144021034	
Epoch: 59 / 100 Training Loss: 135.04605770111084	Test Loss:
36.06994870305061	
Epoch: 60 / 100 Training Loss: 134.73051998019218	Test Loss:
36.03017449378967	
Epoch: 61 / 100 Training Loss: 134.77700686454773	Test Loss:
36.26178798079491	
Epoch: 62 / 100 Training Loss: 134.56750398874283	Test Loss:
36.36930540204048	
Epoch: 63 / 100 Training Loss: 134.3830843269825	Test Loss:
35.89995536208153	
Epoch: 64 / 100 Training Loss: 134.22210678458214	Test Loss:
37.44350093603134	
Epoch: 65 / 100 Training Loss: 134.31491148471832	Test Loss:
36.42756715416908	
Epoch: 66 / 100 Training Loss: 133.98335886001587	Test Loss:
35.9828961789608	
Epoch: 67 / 100 Training Loss: 133.91022527217865	Test Loss:
36.17883452773094	
Epoch: 68 / 100 Training Loss: 133.51893836259842	Test Loss:
35.96130481362343	
Epoch: 69 / 100 Training Loss: 133.90603253245354	Test Loss:
36.07605600357056	
Epoch: 70 / 100 Training Loss: 133.3829669356346	Test Loss:
35.91865962743759	
Epoch: 71 / 100 Training Loss: 133.3340601027012	Test Loss:
36.70209327340126	
Epoch: 72 / 100 Training Loss: 133.73578864336014	Test Loss:
36.085397362709045	
Epoch: 73 / 100 Training Loss: 133.52470207214355	Test Loss:
36.064752370119095	
Epoch: 74 / 100 Training Loss: 133.1799268424511	Test Loss:
36.026156067848206	
Epoch: 75 / 100 Training Loss: 133.26336923241615	Test Loss:
35.974625289440155	
Epoch: 76 / 100 Training Loss: 133.26395693421364	Test Loss:
35.977458477020264	
Epoch: 77 / 100 Training Loss: 132.8995196223259	Test Loss:

36.6982978284359		
Epoch: 78 / 100 Training Loss: 133.10581400990486	Test Loss:	
35.85940346121788		
Epoch: 79 / 100 Training Loss: 132.6943188905716	Test Loss:	
35.95867204666138		
Epoch: 80 / 100 Training Loss: 132.88827127218246	Test Loss:	
35.84911406040192		
Epoch: 81 / 100 Training Loss: 133.28380650281906	Test Loss:	
36.01386970281601		
Epoch: 82 / 100 Training Loss: 132.37537708878517	Test Loss:	
36.04409897327423		
Epoch: 83 / 100 Training Loss: 132.8774455487728	Test Loss:	
35.851144552230835		
Epoch: 84 / 100 Training Loss: 132.53932788968086	Test Loss:	
35.89164027571678		
Epoch: 85 / 100 Training Loss: 132.59623190760612	Test Loss:	
36.02489432692528		
Epoch: 86 / 100 Training Loss: 132.02949604392052	Test Loss:	
36.67078024148941		
Epoch: 87 / 100 Training Loss: 132.04837357997894	Test Loss:	
35.88687840104103		
Epoch: 88 / 100 Training Loss: 132.11943557858467	Test Loss:	
35.89442652463913		
Epoch: 89 / 100 Training Loss: 131.82287102937698	Test Loss:	
35.74980768561363		
Epoch: 90 / 100 Training Loss: 132.0360058248043	Test Loss:	
35.946353524923325		
Epoch: 91 / 100 Training Loss: 131.82789173722267	Test Loss:	
35.85142061114311		
Epoch: 92 / 100 Training Loss: 131.64694225788116	Test Loss:	
35.84598225355148		
Epoch: 93 / 100 Training Loss: 131.85038036108017	Test Loss:	
36.01533231139183		
Epoch: 94 / 100 Training Loss: 131.46432846784592	Test Loss:	
36.32975900173187		
Epoch: 95 / 100 Training Loss: 131.63746419548988	Test Loss:	
36.03915509581566		
Epoch: 96 / 100 Training Loss: 131.45140880346298	Test Loss:	
35.91527062654495		
Epoch: 97 / 100 Training Loss: 131.1892467737198	Test Loss:	
35.95221844315529		
Epoch: 98 / 100 Training Loss: 131.4272505044937	Test Loss:	
36.182045221328735		
Epoch: 99 / 100 Training Loss: 131.37129864096642	Test Loss:	
35.74391394853592		
Epoch: 100 / 100	Training Loss: 131.15149646997452	Test Loss:
36.11794114112854		

```
[45]: rnn_accuracy = accuracy(rnnModel, test_loader3, device)

print(rnn_accuracy)
```

77.955

5.2 Task 5(b): GRU

- Here, we train a GRU on sentence embeddings obtained by stacking the first 10 word embeddings and applying padding if the sentence is smaller than 10 words.
- We train the neural network using AdamW optimizer with 1e-3 learning rate with 1e-2 weight decay for 30 epochs with batch size 256.
- Upon re-running the code several times, we observe the accuracy on test set for this model is roughly around **79.1% - 79.7%**.

```
[46]: class GRU1f(nn.Module):
    def __init__(self):
        super(GRU1f, self).__init__()
        self.embed_size = 300

        self.hidden_size = 10
        self.out_size = 2

        self.gru = nn.GRU(input_size=self.embed_size, hidden_size=self.
↪hidden_size, batch_first=True)
        self.linear = nn.Linear(self.hidden_size, self.out_size)

    def forward(self, x):
        x, hidden = self.gru(x)
        x = self.linear(x[:, -1, :])

        return x

gruModel = GRU1f()
gruModel.to(device)
print(gruModel)
```

```
GRU1f(
  (gru): GRU(300, 10, batch_first=True)
  (linear): Linear(in_features=10, out_features=2, bias=True)
)
```

```
[47]: optimizer = torch.optim.AdamW(gruModel.parameters(), lr=1e-3, weight_decay=1e-2)
```

```
[48]: train(gruModel, train_loader3, test_loader3, criterion, optimizer, n_epochs=30,
↪device=device)
```

Epoch: 1 / 30 Training Loss: 181.33681312203407 Test Loss:
40.29534995555878

Epoch: 2 / 30	Training Loss: 152.9514371752739	Test Loss:
38.26805377006531		
Epoch: 3 / 30	Training Loss: 146.8739361166954	Test Loss:
37.00517484545708		
Epoch: 4 / 30	Training Loss: 143.26274248957634	Test Loss:
36.36405465006828		
Epoch: 5 / 30	Training Loss: 141.10489463806152	Test Loss:
35.926189661026		
Epoch: 6 / 30	Training Loss: 139.08646640181541	Test Loss:
35.67848202586174		
Epoch: 7 / 30	Training Loss: 137.7913582623005	Test Loss:
35.40585646033287		
Epoch: 8 / 30	Training Loss: 136.35473904013634	Test Loss:
35.08504235744476		
Epoch: 9 / 30	Training Loss: 135.20484054088593	Test Loss:
34.93747437000275		
Epoch: 10 / 30	Training Loss: 134.1550863981247	Test Loss:
34.81758573651314		
Epoch: 11 / 30	Training Loss: 133.29672893881798	Test Loss:
34.710423558950424		
Epoch: 12 / 30	Training Loss: 132.54534032940865	Test Loss:
34.62228259444237		
Epoch: 13 / 30	Training Loss: 131.5119285285473	Test Loss:
34.5435888171196		
Epoch: 14 / 30	Training Loss: 131.05372193455696	Test Loss:
34.433901876211166		
Epoch: 15 / 30	Training Loss: 130.31678879261017	Test Loss:
34.589472591876984		
Epoch: 16 / 30	Training Loss: 129.95355436205864	Test Loss:
34.305437207221985		
Epoch: 17 / 30	Training Loss: 128.96858605742455	Test Loss:
34.44904673099518		
Epoch: 18 / 30	Training Loss: 128.68122020363808	Test Loss:
34.321634382009506		
Epoch: 19 / 30	Training Loss: 128.16218599677086	Test Loss:
34.297405660152435		
Epoch: 20 / 30	Training Loss: 127.77931371331215	Test Loss:
34.212056785821915		
Epoch: 21 / 30	Training Loss: 127.33910638093948	Test Loss:
34.577800303697586		
Epoch: 22 / 30	Training Loss: 127.14896404743195	Test Loss:
34.326654225587845		
Epoch: 23 / 30	Training Loss: 126.58092293143272	Test Loss:
34.400550216436386		
Epoch: 24 / 30	Training Loss: 126.18402501940727	Test Loss:
34.0765418112278		
Epoch: 25 / 30	Training Loss: 125.64939358830452	Test Loss:
34.106071442365646		

Epoch: 26 / 30	Training Loss: 125.47544345259666	Test Loss: 34.14732027053833
Epoch: 27 / 30	Training Loss: 125.23535743355751	Test Loss: 34.172821909189224
Epoch: 28 / 30	Training Loss: 124.74122029542923	Test Loss: 33.96618315577507
Epoch: 29 / 30	Training Loss: 124.36532750725746	Test Loss: 34.191488176584244
Epoch: 30 / 30	Training Loss: 124.17429465055466	Test Loss: 34.5347506403923

```
[49]: gru_accuracy = accuracy(gruModel, test_loader3, device)

print(gru_accuracy)
```

79.275

5.3 Task 5(c): LSTM

- Here, we train a LSTM on sentence embeddings obtained by stacking the first 10 word embeddings and applying padding if the sentence is smaller than 10 words.
- We train the neural network using AdamW optimizer with 1e-3 learning rate and 1e-2 weight decay for 30 epochs with batch size 256.
- A relu activation layer between the lstm cell and the linear layer is applied. Also, the output is passed through a tanh layer before being returned by the model.
- Upon re-running the code several times, we observe the accuracy on test set for this model is roughly between **79.2% - 79.8%**.

```
[50]: class LSTMClf(nn.Module):
    def __init__(self):
        super(LSTMClf, self).__init__()
        self.embed_size = 300
        self.hidden_size = 10
        self.out_size = 2

        self.lstm = nn.LSTM(input_size=self.embed_size, hidden_size=self.
↪hidden_size, batch_first=True)
        self.relu = nn.ReLU()
        self.linear = nn.Linear(self.hidden_size, self.out_size)
        self.tanh = nn.Tanh()

    def forward(self, x):
        x, (hidden, cell) = self.lstm(x)
        x = self.relu(x)
        x = self.linear(x[:, -1, :])
        x = self.tanh(x)
```

```

        return x

lstmModel = LSTMClf()
lstmModel.to(device)
print(lstmModel)

```

```

LSTMClf(
  (lstm): LSTM(300, 10, batch_first=True)
  (relu): ReLU()
  (linear): Linear(in_features=10, out_features=2, bias=True)
  (tanh): Tanh()
)

```

```

[51]: optimizer = torch.optim.AdamW(lstmModel.parameters(), lr=1e-3,
    ↪weight_decay=1e-2)

```

```

[52]: train(lstmModel, train_loader3, test_loader3, criterion, optimizer,
    ↪n_epochs=30, device=device)

```

Epoch: 1 / 30	Training Loss: 188.53236678242683	Test Loss: 42.60245108604431
Epoch: 2 / 30	Training Loss: 161.78118962049484	Test Loss: 40.11137869954109
Epoch: 3 / 30	Training Loss: 154.7561023235321	Test Loss: 39.04612699151039
Epoch: 4 / 30	Training Loss: 150.59949985146523	Test Loss: 38.261914163827896
Epoch: 5 / 30	Training Loss: 147.96168661117554	Test Loss: 37.8667289018631
Epoch: 6 / 30	Training Loss: 145.9180660545826	Test Loss: 37.60482919216156
Epoch: 7 / 30	Training Loss: 144.42314419150352	Test Loss: 37.16939628124237
Epoch: 8 / 30	Training Loss: 142.6743516921997	Test Loss: 36.983817517757416
Epoch: 9 / 30	Training Loss: 141.31478962302208	Test Loss: 36.718128740787506
Epoch: 10 / 30	Training Loss: 140.29806298017502	Test Loss: 36.9502349793911
Epoch: 11 / 30	Training Loss: 139.1070382297039	Test Loss: 36.48337149620056
Epoch: 12 / 30	Training Loss: 138.49619647860527	Test Loss: 36.24354547262192
Epoch: 13 / 30	Training Loss: 137.58472111821175	Test Loss: 36.23565372824669
Epoch: 14 / 30	Training Loss: 136.69559437036514	Test Loss: 36.28344339132309
Epoch: 15 / 30	Training Loss: 136.19343376159668	Test Loss:

```

36.044730961322784
Epoch: 16 / 30 Training Loss: 135.28828984498978 Test Loss:
36.178139090538025
Epoch: 17 / 30 Training Loss: 134.70898419618607 Test Loss:
35.870521783828735
Epoch: 18 / 30 Training Loss: 133.9309034049511 Test Loss:
35.881180971860886
Epoch: 19 / 30 Training Loss: 133.92980736494064 Test Loss:
35.94572842121124
Epoch: 20 / 30 Training Loss: 132.89750257134438 Test Loss:
35.94690823554993
Epoch: 21 / 30 Training Loss: 132.76369643211365 Test Loss:
35.764438807964325
Epoch: 22 / 30 Training Loss: 131.8276786506176 Test Loss:
35.75800174474716
Epoch: 23 / 30 Training Loss: 131.56945458054543 Test Loss:
36.201402485370636
Epoch: 24 / 30 Training Loss: 131.06357857584953 Test Loss:
35.585066854953766
Epoch: 25 / 30 Training Loss: 130.93170562386513 Test Loss:
35.64456853270531
Epoch: 26 / 30 Training Loss: 130.27441161870956 Test Loss:
36.144619435071945
Epoch: 27 / 30 Training Loss: 129.86872991919518 Test Loss:
35.68116870522499
Epoch: 28 / 30 Training Loss: 129.67316290736198 Test Loss:
35.78095597028732
Epoch: 29 / 30 Training Loss: 129.42498436570168 Test Loss:
35.652844071388245
Epoch: 30 / 30 Training Loss: 128.7190609574318 Test Loss:
35.79785969853401

```

```

[53]: lstm_accuracy = accuracy(lstmModel, test_loader3, device)

print(lstm_accuracy)

```

79.56